

Optimizations on ECDH25519 and ML-KEM

David Wu

Cryptography Engineering Final Project

December 16, 2025

ML-KEM(Kyber)

- ▶ Selected by NIST as the primary Post-Quantum Cryptography (PQC) standard for general-purpose key encapsulation.
- ▶ Resistant to quantum computer attacks
- ▶ Security is based on the Module Learning With Errors (M-LWE) problem in structured lattices.
- ▶ Operations are performed over polynomial rings $Z_q[x]/x^n + 1$ where $n = 256$ and $q = 3329$

Open Source Implementation

An update on Keccak performance on ARMv7-M

- ▶ Author : Alexandre Adomnicai
- ▶ Faster Keccak Implementation in Assembly with two versions
 - ▶ **Balanced** (My choice)
 - ▶ Fast –Lazy rotations for all rounds which gains 2% performance improvement but with 50% extra code size
- ▶ Grouping memory accesses together as much as possible for pipelining
- ▶ 32-bit interleaving
- ▶ Lazy rotations for 3/4 of the rounds

Optimization Strategies 1 – Polynomial Operations

Small Optimizations

- ▶ Replaces dynamic indexed shifting with incremental shifting to reduce arithmetic overhead within **cbd2**.
- ▶ Implement the whole base multiplication directly to reduce calling overhead

DSP Instructions

- ▶ **poly_add** and **poly_sub** is suitable for DSP operations.
- ▶ **uadd16** / **usub16** can reach 2x speedup for these functions.

Compression / Decompression

- ▶ Intermediate bit-packing entirely within registers to eliminate redundant memory access
- ▶ Barrel Shifter to embed bitwise shifts into logical instructions
- ▶ Merging scaling factors into multiplication constants with post-increment addressing for pointer updates.

Optimization Strategies 2 – NTT

NTT Optimization

- ▶ Performance Bottleneck : Load / Store Operations
 - ▶ Loop Unrolling
 - ▶ Layer Fusing – $7 = 2 + 2 + 2 + 1$
 - ▶ Inverse NTT with Cooley-Tukey Butterfly to achieve lazy reduction
 - ▶ Merge twist factors and scaling factors for each element
- ▶ Performance Improvement : Two at a time
 - ▶ The elements for NTT are 16-bit, but the registers are 32-bit.
 - ▶ **ldr / str** to load and store neighboring elements together
 - ▶ DSP instructions to achieve SIMD arithmetic operations –
uadd16 / usub16
 - ▶ Specialized Montgomery Reduction for two elements from Faster Kyber and Dilithium on the Cortex-M4
 - ▶ Authors : Amin Abdulrahman / Vincent Hwang / Matthias J. Kannwischer / Amber Sprengels

Optimization Results 1 – ML-KEM

C Reference Implementation

- ▶ **poly_ntt** : 35830 Cycles
- ▶ **Keypair Generation** : 1325464 Cycles
- ▶ **Encapsulation** : 1481614 Cycles
- ▶ **Decapsulation** : 1741181 Cycles

Until 12 / 15

- ▶ **poly_ntt** : 19746 Cycles (44.89% Speedup)
- ▶ **Keypair Generation** : 772394 Cycles(41.73% Speedup)
- ▶ **Encapsulation** : 823690 Cycles(44.41% Speedup)
- ▶ **Decapsulation** : 960378 Cycles(44.84% Speedup)

Optimization Results 2 – Specific Functions and Code Size

Individual Functions

- ▶ **NTT** : 23546 → 14566 (38.1% Reduction)
- ▶ **Inverse NTT** : 32376 → 15469 (52.2% Reduction)
- ▶ **poly_add** : 2056 → 1045 (49.2% Reduction)
- ▶ **poly_sub** : 2343 → 1045 (55.4% Reduction)

NTT code size in `ntt.S`

- ▶ Overall 2028 bytes
- ▶ However, it includes base multiplication + NTT + Inverse NTT
- ▶ Each of them \leq 1024 Bytes

Next Steps

- ▶ Find possibilities to employ FPRs to reduce memory operations
- ▶ More aggressive layer fusing in NTT within code size constraint
- ▶ Implement functions with loops in assembly to reduce calling overhead with possible unrolling
- ▶ Optimizing redundant **mov**
- ▶ Register allocation and pipelining

- ▶ An Elliptic Curve Diffie-Hellman key exchange protocol using Curve25519.
- ▶ Allows two parties to securely agree on a shared secret over an insecure channel.
- ▶ 128-bit security level with prime $p = 2^{255} - 19$, allowing fast arithmetic on CPUs.

Let G be the standard base point on the curve.

1. **Alice:** Generates private a , sends public $A = a \cdot G$.
2. **Bob:** Generates private b , sends public $B = b \cdot G$.
3. **Shared Secret:** Both compute $S = a \cdot B = b \cdot A = (a \cdot b) \cdot G$.

Optimization Strategies 1

Basic Side-Channel Prevention

- ▶ Constant time double-and-add instead of variable scalar multiplication with secret-dependent branch
- ▶ Branchless cmov and cswap helper

Finite Field Arithmetic

- ▶ Radix- 2^8 for all operations → Radix $2^{25.5}$
 - ▶ Before : 1024 Multiplications and additions without overflow risk. However, register is 32-bit wide.
 - ▶ Now : Repacking without freezing in the middle of multiplication(squaring). Only about 100 multiplications needed. (90% improvement)
- ▶ Multiplication and addition → **smlal** and **smull** in assembly
- ▶ Nested Loop Operand Scanning → Product Scanning
- ▶ Specialized Squaring to reduce repeated multiplications(results later)

Optimization Strategies 2

Scalar Multiplication

- ▶ Fixed-base Window Method with precomputed table($w = 4$)
- ▶ Conditional move to achieve constant-time loading
- ▶ Calculate base $G, 2^w G, 2^{2w} G, \dots$
- ▶ $table[i][j] = j \times 2^{iw} G$
- ▶ Reduce 255 doubling + 128 additions to 64 additions only
- ▶ Large size : 512 KB (within 1MB flash)

Optimization Results – ECDH25519

C Reference with constant-time Implementation

- ▶ **crypto_scalarmult_base** : 54552132 Cycles
- ▶ **crypto_scalarmult** : 51772998 Cycles

Until 12 / 15

- ▶ **crypto_scalarmult_base** : 3141710 Cycles (94.24% Speedup)
- ▶ **crypto_scalarmult** : 12454872 Cycles (75.94% Speedup)

More Specifically...

- ▶ **fe25519_mul** : 9695045 → 1812045 Cycles
(81.31% Speedup)
- ▶ **fe25519_sqr** : 9695078 → 1431045 Cycles (85.24% Speedup)

Next Step

- ▶ Larger radix in other functions
- ▶ FPRs usage to reduce stack usage
- ▶ Pipelining and instruction selection
- ▶ Be aware of secret-dependent operations

References I

- [Abd+22] Amin Abdulrahman, Vincent Hwang,
Matthias J. Kannwischer, and Amber Sprenkels. *Faster
Kyber and Dilithium on the Cortex-M4*. Cryptology
ePrint Archive, Paper 2022/112. 2022. URL:
<https://eprint.iacr.org/2022/112>.
- [Ado23] Alexandre Adomnicai. *An update on Keccak
performance on ARMv7-M*. Cryptology ePrint Archive,
Paper 2023/773. 2023. URL:
<https://eprint.iacr.org/2023/773>.